




# On-demand Research Compute Environments

# Flight Compute Introduction - Agenda



- Introductions
- Deploying your cluster
  - Requirements and configuration options
  - Deploying with CloudFormation templates
- Using your cluster
  - Your user environment and module files
  - Start and managing sessions
  - Applications
  - Using file and object storage
  - Job-scheduler; submitting and running jobs
- Terminating your cluster



**alces**flight

# Alces Flight

## Overview

- Research compute environment; including:
  - Familiar Linux HPC cluster look and feel
  - HPC applications, libraries, MPIs and tools
  - Batch scheduler
  - Data management and accessibility tools
- Target audience
  - New and existing users
  - Scientific researchers and end-users
  - Ephemeral usage; rapid start and stop
- Unique selling points
  - Minimum possible setup time
  - Extremely customisable and expandable
  - Leverages benefits of public cloud



- Flight is an enabler for HPC on AWS:
  - Handles compute node start/stop efficiently, reducing idle compute cycles
  - Applications and data are centralised, minimising costs while cluster is idle
  - Simple scaling of clusters via batch-scheduler job submission
  - Ephemeral usage encourages users to manage their data better
  - Large library of available applications optimised for available instance types
  - Users can access from anywhere - no site requirements
  - Groups of users can collaborate on a single platform from any location
  - Standardised platform for training users new to AWS and HPC



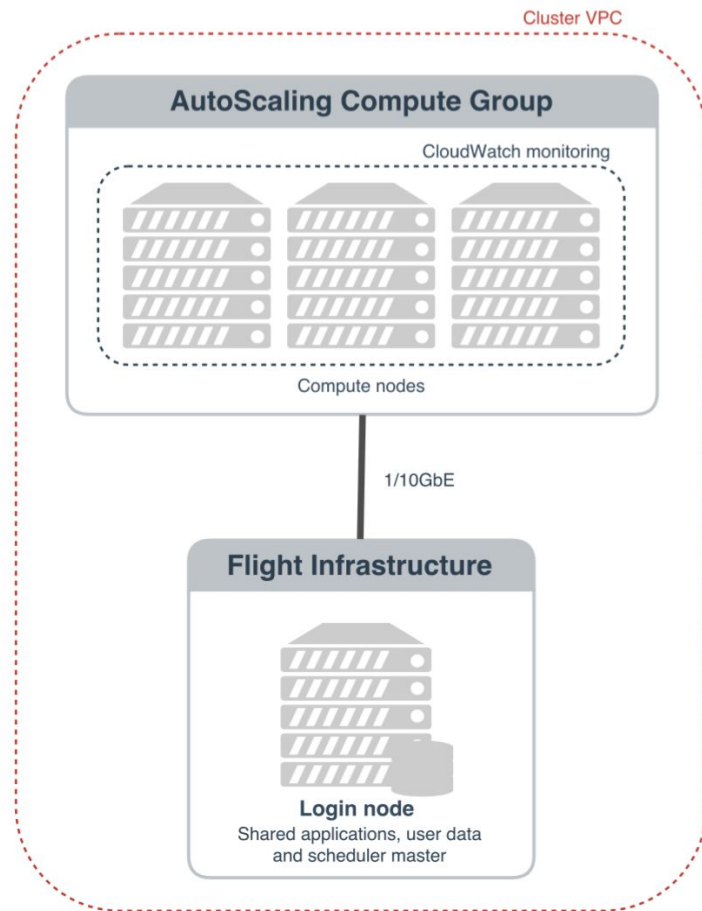
- Flight Compute
  - Launched from Marketplace: Single AMI with CloudFormation template
  - One login node + one or more compute nodes
  - Auto-configured shared NFS home directories
  - Batch job scheduler with example job-scripts
  - Access to more than 750 applications, + session and data management tools
- Additional support appliances available soon
  - Flight Access Manager
  - Flight Application Manager
  - Flight Storage Manager
  - Flight Galaxy appliance for BioInformatics pipelines

# Deploying your cluster

# Deploying your cluster



- Requirements for launching
  - *Marketplace* or CFN template launch
  - SSH client and SSH keypair registered in AWS
  - AWS user account with IAM privileges
- Deployed architecture
  - Single cluster master node, hosting:
    - Cluster job scheduler,
    - Shared applications and module files
    - NFS area for user data
  - AutoScaling compute node group
    - Auto-detect local master and auto-configure
  - New VPC to contain the cluster



# How do I deploy?



- Deployment from AWS Marketplace using CloudFormation
- Search for “**Alces Flight**” or visit <http://tiny.cc/alcesflight>
- Templates also available to public from S3 URL
  - <https://s3.amazonaws.com/alces-flight-templates/2016.2r3/demo.json>
- The templates create:
  - VPC, subnet, gateway, security group
  - A single cluster master node instance with EBS storage (sg2 for *demo*)
  - Choice of on-demand or SPOT instances in an auto-scaling group
  - CloudWatch rule created along with IAM role to report data to AWS
  - *Demo* template enables all software application repos



# Other Launch options



- AWS Marketplace
  - Simple method for generating a ephemeral, personal cluster
- Alces templates
  - Common options for specific purposes (e.g. fixed-size, demo, etc.)
- Manual launch
  - Create an instance using the AMI and run the configuration tool
- User-created templates
  - Start from Marketplace template
  - Customise to deliver specific solutions



# Deployment from Marketplace - steps

- Select the **Continue** button from the AWS Marketplace product page
- Select the **Custom Launch** tab from the Marketplace launch page
- Select the AWS region in which you'd like your cluster to run
- Select the **Personal HPC Compute Cluster** option
- Then click the **Launch with CloudFormation console** button to start
- In the console, click **Next** to begin configuring your environment



- With your own preferences, configure the following options
  - **Stack name**; the name for your cluster
  - **ComputeAutoscaling**; enter 0 to disable or 1 to enable auto-scaling
  - **ComputeSpotPrice**; spot price to be used; enter 0 for on-demand
  - **ComputeType**; select your compute-node instance type
  - **FlightCustomBucket**; enter an S3 bucket containing customisation information for your cluster.
  - **FlightCustomProfiles**; enter the names of the customisation profiles to use, separated by spaces.
  - **InitialNodes**; enter the number of nodes to start immediately in this box in an auto-scaling cluster.
  - **Keypair**; choose an existing AWS keypair to launch your Flight cluster with.
  - **LoginSystemDiskSize**; choose the size of your login node disk - the shared filesystem for your cluster.
  - **LoginType**; use the drop-down box to choose the AWS instance type for your login node.
  - **MaxNodes**; enter the maximum size that your cluster will scale to, up to a maximum of 32 nodes.
  - **NetworkCIDR**; enter a network range permitted to access your cluster; to open to all, enter “0.0.0.0/0”
  - **Username**; enter the username you want to use to connect to the cluster.



- From the **Options** page - click the **Next** button to review your settings
- From the **Review** page - click on the checkbox in the **Capabilities** section to allow creation of IAM role, then click the **Create** button to launch your cluster
- The cluster will begin creating its resources.
- Once the environment has finished creating, the public IP address of the cluster login node will be available from the **Outputs** tab

| Overview | Outputs       | Resources | Events  | Template | Parameters | Tags | Stack Policy | Change Sets |
|----------|---------------|-----------|---|----------|------------|------|--------------|-------------|
| Key      | Value         |           | Description   |          |            |      |              |             |
| AccessIP | 52.50.133.133 |           | Public access IP for your Flight Compute environment. ... |          |            |      |              |             |

- Using your AWS key - SSH to the public IP address user you specified

# Using the Compute Cluster

Example use-cases and workloads using your Alces Flight environment

# Linux User Environment

## Useful commands:

```
qhost  
pwd  
cd, mkdir, rmdir  
cp, mv, rm  
ssh ip-10-75-0-187  
yum search <package>  
sudo yum install <package>
```

- Your user account has full *sudo* access
  - Become root user as required
- Password-less SSH is automatically enabled between nodes
  - Use the **qhost** command to view cluster nodes
  - Log in to other nodes using SSH
- Shared storage for user data
  - Your home-directory is mounted on all nodes
  - Home directory is **/home/<username>/** or **~/**
    - e.g. **/home/alces/**
- Standard Linux environment
  - Install new packages on a single node using **yum**:
    - e.g. **sudo yum install nano**
  - Linux environment variables:
    - **\$PATH**
    - **\$LD\_LIBRARY\_PATH**
    - **\$MANPATH**
    - **\$USER**
    - **\$PWD**

# Working with multiple nodes

## Useful commands:

`nodeattr`

`pdsh`

- Automatic node grouping support (genders)
  - Nodes are automatically added/removed
    - Default group is called **nodes**
  - Scales with compute cluster
  - Use **nodeattr** command to enumerate
    - e.g `nodeattr -s nodes`
- Parallel Distributed Shell (PDSH)
  - Use the **pdsh** command to run commands on multiple nodes
  - Manually specify nodes, or use genders support
  - Syntax:

```
pdsh [-g group,group,... | -w host,host,...] command
```

- Examples:

```
pdsh -g nodes uptime
```

```
pdsh -g nodes -x ip-10-75-0-123 'df -h /tmp'
```

```
pdsh -w ip-10-75-0-[123,125] 'tail -f /tmp/myfile.out'
```

```
pdsh -g nodes 'sudo yum -y install nano'
```

```
pdsh -g nodes -f 1 date
```

# Graphical Desktop

## Useful commands:

```
alces session avail
alces session enable <type>
alces session start <type>
alces session list
alces session info <id>
alces session kill <id>
xrandr
xrandr -s 1280x1024
```

- Graphical sessions can be launched on the login node
- Use the **alces session avail** command to list available types
- Use **alces session enable** command to enable a new type
  - e.g. `alces gridware enable chrome`
- Create a new session with **alces session start** command
  - e.g. `alces gridware start gnome`
- VNC session will be started with connection info printed:

```
[alces@login1(mycluster) ~]$ alces session start gnome
VNC server started:
  Identity: 5650cbf4-04c0-11e6-95a3-0af76f7bf051
  Type: gnome
  Host: 52.50.133.113
  Port: 5901
  Display: 1
  Password: u4NXvjpy
```

Depending on your client, you can connect to the session using:

```
vnc://alces:u4NXvjpy@52.50.133.133:5901
52.50.133.133:5901
52.50.133.133:1
```

MacOSX URL

VNC client URL

- Use a VNC client to connect using the URL provided



# Environment Modules

## Useful commands:

```
module help
```

- Environment modules help users find and run applications, libraries and compilers
- Modules configure your shell variables
  - `$PATH`, `$LD_LIBRARY_PATH`, `$MANPATH`
  - For the user's current running session only (by default)
- Can be run at the command-prompt or in job-scripts
- New apps are installed with supporting module files
- Examples:

```
module avail
module load libs/gcc compilers/gcc
module list
module unload compilers/gcc
module whatis services/gridscheduler
module keyword gcc
module purge
module initadd apps/bowtie2
```

# Cluster Software

## Useful commands:

```
module avail  
alces gridware list
```

- Alces Gridware project
  - Holds instructions for installing applications
  - Includes libraries, compilers and MPIs
  - Centrally manages applications on the login node
- Users can use gridware to
  - Download, compile and install software
  - Install depots of pre-built software
  - Prepare the environment to run jobs
- Applications are installed with a module file
- Use **module avail** command to list installed applications
- Example **alces gridware** commands:

```
alces gridware list  
alces gridware search --names gromacs  
alces gridware info apps/gromacs/4.6.1
```

# Installing Applications

## Useful commands:

```
alces gridware list
alces gridware install
  apps/bowtie2
module whatis apps/bowtie2
alces gridware purge
  apps/bowtie2
```

- New applications installed using **alces gridware install**
  - Specify the version required
  - Specify any variants (if required)

e.g. `alces gridware install apps/bowtie2/2.2.6`

- Use **module avail** to view new module file once installed
- Use **module whatis apps/bowtie2** to view license
- Remove apps using **module purge apps/myapp**
- Multiple versions of applications can be installed
- Applications are available across all nodes
- Applications are installed into **local** depot

# Gridware Depots

## Useful commands:

```
alces gridware depot list
alces gridware depot fetch
alces gridware depot enable
module avail
```

- Applications can be collected together in depots
- Default location for new apps is **local** depot
- Depots are pre-compiled for AWS
- Install a depot using **alces gridware depot install**

e.g. `alces gridware depot install development`

- Use **alces gridware depot list** to view installed depots
- Use **module avail** to view new applications
- Depots can be used to:
  - Group applications together
  - Handle binary/commercial apps
  - Export applications for quick access

# Volatile Repository

## Useful commands:

```
alces gridware list
alces gridware update
nano
/opt/gridware/etc/gridware.yml
```

- Applications are stored in different repos
- **Main** repo
  - Binary versions of applications available
  - Compatible with auto-scaling clusters
- **Volatile** repo
  - Compiled from source on demand
  - Requires static clusters
- Marketplace product has **main** repo enabled by default
- Enable volatile repo by editing:
  - `/opt/gridware/etc/gridware.yml`
  - Remove comment “#” in repo config for volatile

# Storing Data Files

## Useful commands:

```
scp -i key <file> <remote>
```

```
sftp <file> <remotefile>
```

```
rsync -ravP -e "ssh -i key"  
  <localfile> <remotefile>
```

- Users can copy files to the cluster using SCP
- Uses SSH keypair for passwordless transfer
- Data is encrypted in transit
- Many desktop clients available:
  - scp/sftp (Linux, Mac)
  - pscp/psftp (Windows)
  - WinSCP, FileZila (Windows explorer interface)
- Specify keypair, username and path to files:

e.g. `scp -i aws.pub alces@52.50.133.133:/home/alces/file .`

- Tools like **rsync** can make copying faster by skipping files that haven't changed:

e.g. `rsync -ravP -e "ssh -i aws.pub" /storage/data  
alces@52.50.133.133:/home/alces/.`

# Archiving Data

## Useful commands:

```
alces storage enable
alces storage configure
alces storage avail
alces storage use
alces storage get
alces storage put
```

- The **alces storage** command allows users to store files in S3, dropbox and POSIX services
- Use the **alces storage enable** command:
  - **s3** - Amazon S3 object storage
  - **dropbox** - Dropbox object storage
  - **posix** - An existing POSIX path
- Use the **alces storage configure** command to set up a new storage service

e.g. `alces storage configure mys3 s3`

- Select a service with **avail** and **use** commands
- Upload and download files with **get** and **put**

e.g. `alces storage put localfile s3://remotefile`  
`alces storage -r get s3://jobs .`  
`alces storage mkbucket newbucket`

# Running an Application on the login node

## Useful commands:

```
alces session start  
xrandr  
module load
```

- Batch jobs are submitted to run on compute nodes via the job scheduler
- Users can also run interactively on the login node
  - Example work-flow:

```
alces session start gnome
```

<connect to VNC desktop>

<from the desktop **Application** menu, select **Terminal**>

```
glxgears
```



# Using the Job Scheduler

- Users submit jobs to a batch queue which are then run on compute nodes
- The job scheduler is responsible for:
  - Starting and stopping jobs
  - Assembling machinefile for parallel jobs
  - Storing printed job output in a file
- Users provide a **job script** which instructs the scheduler what to do to run their application
- Users can use **directives** to instruct the scheduler how to run their job; these may be
  - Included in the user's job-script
  - Provided on the command-line at submission time

# Writing a job script

## Useful commands:

```
qstat
```

```
qsub <job-script name>
```

```
cat <output file>
```

- A simple job-script:

```
#!/bin/bash -l
echo Hello from job $JOB_ID
sleep 60
echo I was run on node $HOSTNAME
```

- Submit the job using **qsub <jobscrip>**
  - View the job status using **qstat**
    - Output shows your unique job ID
    - Job status (r=running, qw=queuing/waiting)
  - By default, job output stored in home-dir
    - Filename is **~/<jobscrip-name>.o<job-ID>**
- e.g. `cat ~/myjob.sh.o5`

# Job scheduler directives

## Useful commands:

```
qstat
```

```
qsub <job-script name>
```

```
cat <output file>
```

- The directive “**-o <filename>**” instructs the scheduler where to store job output
  - The directive can be added at submission time:

e.g. `qsub -o /home/alces/myoutput myjob.sh`

- The directive can also be added in the job-script:

```
#!/bin/bash -l
#$ -o /home/alces/myoutput
echo Hello from job $JOB_ID
sleep 60
echo I was run on node $HOSTNAME
```

- Lines in a job-script starting with `#$` are evaluated at submission time

# Job scheduler directives

## Useful commands:

`qstat`

`qsub <directive> <job-script name>`

- More scheduler directives:

|   |   |
|---|---|
| <code>-o &lt;filename&gt;</code>              | Set job output file location  |
| <code>-hold &lt;job-ID,...&gt;</code>         | Wait until the specified job IDs have completed running                 |
| <code>-pe &lt;PEname&gt; &lt;slots&gt;</code> | Use named parallel environment, requesting a number of slots            |
| <code>-p &lt;priority&gt;</code>              | Set the job priority (0=normal, -1000=low)                              |
| <code>-r &lt;yes no&gt;</code>                | Set the job to automatically rerun if terminated                        |
| <code>-l h_vmem=16G</code>                    | Request an amount of usable memory per slot (defaults to total / slots) |
| <code>-l rt=36:00:00</code>                   | Sets max job runtime to 36 hours  |

# Example: Running an MPI job

Useful commands:

```
qsub imb.sh
```

- The **Intel Messaging Benchmark** (IMB) tests latency and bandwidth between compute nodes
- Write a simple job script that launches the IMB application, and submit to the cluster scheduler
- Watch the output to view results

```
#!/bin/bash -l
#$ -j y -N imb
#$ -o $HOME/imb.$JOB_ID.out
#$ -pe mpinodes-verbose 2 -cwd -V

module load apps/imb
mpirun -np 2 -npernode 1 IMB-MPI1
```

# Template job-scripts

## Useful commands:

```
alces template list
```

```
alces template copy <name>  
    <new jobscript>
```

```
qsub <jobscript>
```

- A reference set of job-scripts are available using the **alces template** command
  - **alces template list** shows a list of available templates
    - **simple** - single-core serial job
    - **simple-array** - multiple single-core serial jobs
    - **smp** - multi-threaded job on a single node
    - **mpi-slots** - parallel job across multiple nodes, allocated by number of CPU cores
    - **mpi-nodes** - parallel job across multiple nodes, allocated by number of nodes
- The **alces template copy** command allows users to create new job scripts using a template

```
[alces@login1(tiny) jobs]$ alces template copy smp  
mysmpjob.sh  
alces template copy: template 'smp' copied to  
'mysmpjob.sh'
```

# Managing Jobs

## Useful commands:

`qstat -j <job-ID>`

`qacct -j <job-ID>`

`qsub <job-script name>`

`cat <output file>`

- Users can view the job queue using the **qstat** command; use **qstat -f** for full output
- Individual jobs can be deleted using **qdel <job-ID>**
- Job resource usage is shown using the **qacct <job-ID>** command (once job is completed)
- The status of compute nodes is shown using the **qhost** command

```
[alces-cluster@vlogin1(tiny) jobs]$ qhost
```

| HOSTNAME | ARCH      | NCPU  | LOAD  | MEMTOT | MEMUSE | SWAPTO | SWAPUS |
|----------|-----------|-------|-------|--------|--------|--------|--------|
| -----    | -----     | ----- | ----- | -----  | -----  | -----  | -----  |
| global   | -         | -     | -     | -      | -      | -      | -      |
| vnnode01 | linux-x64 | 20    | 0.01  | 58.8G  | 620.8M | 32.0M  | 0.0    |
| vnnode02 | linux-x64 | 20    | 0.02  | 58.8G  | 638.0M | 32.0M  | 0.0    |
| vnnode03 | linux-x64 | 20    | 0.01  | 58.8G  | 651.0M | 32.0M  | 0.0    |
| vnnode04 | linux-x64 | 20    | 0.01  | 58.8G  | 622.5M | 32.0M  | 0.0    |
| vnnode05 | linux-x64 | 20    | 0.01  | 58.8G  | 682.0M | 32.0M  | 0.0    |
| vnnode06 | linux-x64 | 20    | 0.01  | 58.8G  | 642.6M | 32.0M  | 0.0    |
| vnnode07 | linux-x64 | 20    | 0.01  | 58.8G  | 667.6M | 32.0M  | 0.0    |
| vnnode08 | linux-x64 | 20    | 0.01  | 58.8G  | 642.5M | 32.0M  | 0.0    |
| vnnode09 | linux-x64 | 20    | 0.01  | 58.8G  | 661.3M | 32.0M  | 0.0    |



# Terminating your cluster

- Before termination, store any job-scripts, data and output files created
  - Use **alces storage put -r** command to check-in data you want to keep
- List any applications you've installed
  - Use **module list** to view applications installed
  - Use **alces gridware depot export** to save your applications as a TAR file
- Delete the cluster stack from the *AWS CloudFormation* console
- Once deleted, users are no longer billed for the instances or software





# Start computing today

- Many applications to explore
- Lots more help available online
  - <http://docs.alces-flight.com>



**alces**flight

*Thank you*

[www.alces-flight.com](http://www.alces-flight.com)

<http://alces-flight.github.io>