




On-demand Research Compute Environments



- Introductions
- Deploying your cluster
 - Requirements and configuration options
 - Deploying with CloudFormation templates
- Using your cluster
 - Your user environment and module files
 - Start and managing sessions
 - Applications
 - Using file and object storage
 - Job-scheduler; submitting and running jobs
- Terminating your cluster



alcesflight

Alces Flight

Overview

- Research compute environment; including:
 - Familiar Linux HPC cluster look and feel
 - HPC applications, libraries, MPIs and tools
 - Batch scheduler
 - Data management and accessibility tools
- Target audience
 - New and existing users
 - Scientific researchers and end-users
 - Ephemeral usage; rapid start and stop
- Unique selling points
 - Minimum possible setup time
 - Extremely customisable and expandable
 - Leverages benefits of public cloud



- Flight is an enabler for HPC on AWS:
 - Handles compute node start/stop efficiently, reducing idle compute cycles
 - Applications and data are centralised, minimising costs while cluster is idle
 - Simple scaling of clusters via batch-scheduler job submission
 - Ephemeral usage encourages users to manage their data better
 - Large library of available applications optimised for available instance types
 - Users can access from anywhere - no site requirements
 - Groups of users can collaborate on a single platform from any location
 - Standardised platform for training users new to AWS and HPC



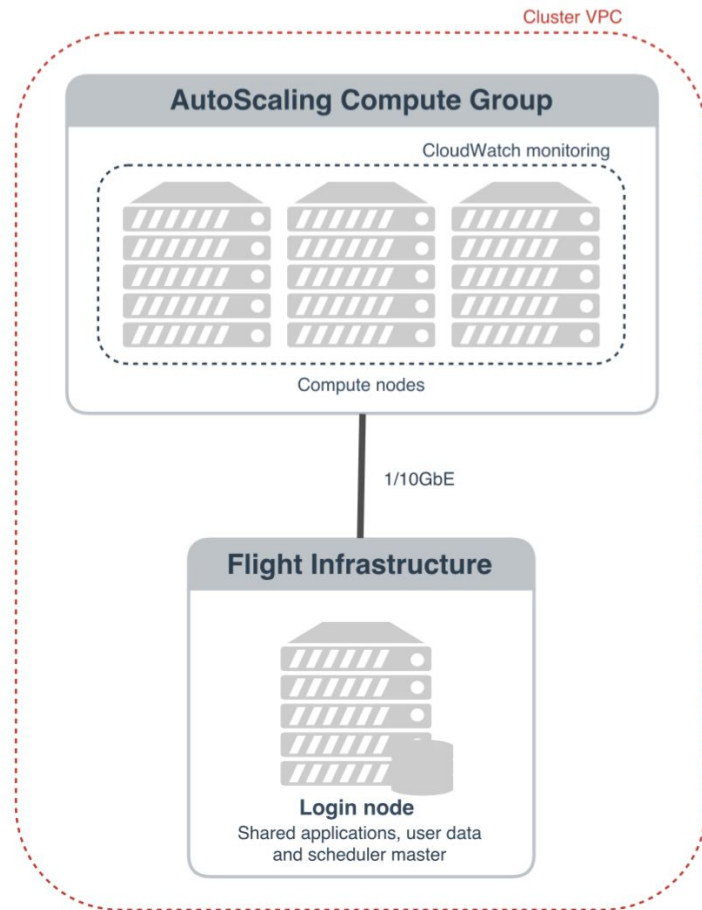
- Flight Compute
 - Launched from Marketplace: Single AMI with CloudFormation template
 - One login node + one or more compute nodes
 - Auto-configured shared NFS home directories
 - Batch job scheduler with example job-scripts
 - Access to more than 750 applications, + session and data management tools
- Additional support appliances available soon
 - Flight Access Manager
 - Flight Application Manager
 - Flight Storage Manager
 - Flight Galaxy appliance for BioInformatics pipelines

Deploying your cluster

Deploying your cluster



- Requirements for launching
 - *EU-West/Ireland region selected (pre-MP)*
 - SSH client and SSH keypair registered in AWS
 - AWS user account with IAM privileges
- Deployed architecture
 - Single cluster master node, hosting:
 - Cluster job scheduler,
 - Shared applications and module files
 - NFS area for user data
 - AutoScaling compute node group
 - Auto-detect local master and auto-configure
 - New VPC to contain the cluster



How do I deploy?



- Deployment using AWS CloudFormation
- Templates available to public and soon through AWS marketplace
 - <https://s3-eu-west-1.amazonaws.com/flight-aws-marketplace/2016.1/8-node.json>
 - <https://s3-eu-west-1.amazonaws.com/flight-aws-marketplace/2016.1/16-node.json>
 - <https://s3-eu-west-1.amazonaws.com/flight-aws-marketplace/2016.1/32-node.json>
- The templates create:
 - VPC, subnet, gateway, security group
 - A single cluster master node instance with EBS storage
 - Initial fixed number of compute nodes
 - Deployed as SPOT instances in an auto-scaling placement group
 - CloudWatch rule created along with IAM role to report data to AWS

Template options



- Three templates initially available in Marketplace
- 8-node template
 - Initial 8 dedicated compute hosts using **c4.large** instance type - 2 cores and 4GB RAM
 - Recommended for evaluation and test usage
- 16-node template
 - Initial 16 dedicated compute hosts using **c3.4xlarge** instance type - 16 cores and 30GB RAM
 - Two SSDs per compute host for local scratch storage
 - Recommended for research use
- 32-node template
 - Initial 32 dedicated compute hosts using **c4.8xlarge** instance type - 36 cores and 60GB RAM
 - Recommended for research use



Deployment - steps

- From the EC2 management dashboard, navigate to the CloudFormation console
- Click **Create New Stack**
- Click **Specify an S3 template URL** and enter one of the previously specified template URLs
- Click **Next** to begin configuring your environment



- With your own preferences, configure the following options
 - **Stack name:** Enter a name for your cluster, e.g. **hpc1**
 - **KeyPair:** Select your AWS keypair
 - **LoginType:** This defines the login node type - setting the number of cores and memory available to the login instance. T-instance types have an 8GB disk; C-instances types have 100GB.
 - **NetworkCIDR:** For increased security, limit access to the login node to this subnet
 - **SpotPrice:** Enter the price in USD per hour for the deployed spot instances, note - this is the maximum cost per hour you are prepared to pay
 - **SystemDiskSize:** Set the root disk size to use for each compute instance.
 - **Username:** Set the user login name that your AWS keypair will be associated with
- Click the **Next** button once you have verified your settings



- From the **Options** page - click the **Next** button to review your settings
- From the **Review** page - click on the checkbox in the **Capabilities** section to allow creation of IAM role, then click the **Create** button to launch your cluster
- The cluster will begin creating its resources.
- Once the environment has finished creating, the public IP address of the cluster login node will be available from the **Outputs** tab

Overview	Outputs	Resources	Events	Template	Parameters	Tags	Stack Policy	Change Sets
Key		Value		Description				
AccessIP		52.50.133.133		Public access IP for your Flight Compute environment. ...				

- Using your AWS key - SSH to the public IP address user you specified

Using the Compute Cluster

Example use-cases and workloads using your Alces Flight environment

Linux User Environment

Useful commands:

```
qhost
pwd
cd, mkdir, rmdir
cp, mv, rm
ssh ip-10-75-0-187
yum search <package>
sudo yum install <package>
```

- Your user account has full *sudo* access
 - Become root user as required
- Password-less SSH is automatically enabled between nodes
 - Use the **qhost** command to view cluster nodes
 - Log in to other nodes using SSH
- Shared storage for user data
 - Your home-directory is mounted on all nodes
 - Home directory is **/home/<username>/** or **~/**
 - e.g. **/users/alces/**
- Standard Linux environment
 - Install new packages on a single node using **yum**:
 - e.g. **sudo yum install nano**
 - Linux environment variables:
 - **\$PATH**
 - **\$LD_LIBRARY_PATH**
 - **\$MANPATH**
 - **\$USER**
 - **\$PWD**

Working with multiple nodes

Useful commands:

`nodeattr`

`pdsh`

- Automatic node grouping support (genders)
 - Nodes are automatically added/removed
 - Default group is called **nodes**
 - Scales with compute cluster
 - Use **nodeattr** command to enumerate
 - e.g `nodeattr -s nodes`
- Parallel Distributed Shell (PDSH)
 - Use the **pdsh** command to run commands on multiple nodes
 - Manually specify nodes, or use genders support
 - Syntax:

```
pdsh [-g group,group,... | -w host,host,...] command
```

- Examples:

```
pdsh -g nodes uptime
```

```
pdsh -g nodes -x ip-10-75-0-123 'df -h /tmp'
```

```
pdsh -w ip-10-75-0-[123,125] 'tail -f /tmp/myfile.out'
```

```
pdsh -g nodes 'sudo yum -y install nano'
```

```
pdsh -g nodes -f 1 date
```

Graphical Desktop

Useful commands:

```
alces session avail
alces session enable <type>
alces session start <type>
alces session list
alces session info <id>
alces session kill <id>
xrandr
xrandr -s 1280x1024
```

- Graphical sessions can be launched on the login node
- Use the **alces session avail** command to list available types
- Use **alces session enable** command to enable a new type
 - e.g. `alces gridware enable chrome`
- Create a new session with **alces session start** command
 - e.g. `alces gridware start gnome`
- VNC session will be started with connection info printed:

```
[alces@login1(mycluster) ~]$ alces session start gnome
VNC server started:
  Identity: 5650cbf4-04c0-11e6-95a3-0af76f7bf051
  Type: gnome
  Host: 52.50.133.113
  Port: 5901
  Display: 1
  Password: u4NXvjpY
```

Depending on your client, you can connect to the session using:

```
vnc://alces:u4NXvjpY@52.50.133.133:5901
52.50.133.133:5901
52.50.133.133:1
```

MacOSX URL

VNC client URL

- Use a VNC client to connect using the URL provided

Environment Modules

Useful commands:

```
module help
```

- Environment modules help users find and run applications, libraries and compilers
- Modules configure your shell variables
 - `$PATH`, `$LD_LIBRARY_PATH`, `$MANPATH`
 - For the user's current running session only (by default)
- Can be run at the command-prompt or in job-scripts
- New apps are installed with supporting module files
- Examples:

```
module avail
module load libs/gcc compilers/gcc
module list
module unload compilers/gcc
module whatis services/gridscheduler
module keyword gcc
module purge
module initadd apps/bowtie2
```

Cluster Software

Useful commands:

```
module avail  
alces gridware list
```

- Alces Gridware project
 - Holds instructions for installing applications
 - Includes libraries, compilers and MPIs
 - Centrally manages applications on the login node
- Users can use gridware to
 - Download, compile and install software
 - Install depots of pre-built software
 - Prepare the environment to run jobs
- Applications are installed with a module file
- Use **module avail** command to list installed applications
- Example **alces gridware** commands:

```
alces gridware list  
alces gridware search --names gromacs  
alces gridware info apps/gromacs/4.6.1
```

Installing Applications

Useful commands:

```
alces gridware list
alces gridware install
  apps/bowtie2
module whatis apps/bowtie2
alces gridware purge
  apps/bowtie2
```

- New applications installed using **alces gridware install**
 - Specify the version required
 - Specify any variants (if required)

e.g. `alces gridware install apps/bowtie2/2.2.6`

- Use **module avail** to view new module file once installed
- Use **module whatis apps/bowtie2** to view license
- Remove apps using **module purge apps/myapp**
- Multiple versions of applications can be installed
- Applications are available across all nodes
- Applications are installed into **local** depot

Gridware Depots

Useful commands:

```
alces gridware depot list
alces gridware depot fetch
alces gridware depot enable
module avail
```

- Applications can be collected together in depots
- Depots are pre-compiled for AWS
- Install a depot using **alces gridware depot fetch**

```
e.g. sudo su - -l -c 'alces gridware depot fetch
https://s3-eu-west-1.amazonaws.com/packages.alces-soft
ware.com/depots/demo'
```

- Use **alces gridware depot list** to view installed depots
- Use **alces gridware depot enable** to enable new depots for a user

```
e.g. alces gridware depot enable demo
```

- Use **module avail** to view new applications

Storing Data Files

Useful commands:

```
scp -i key <file> <remote>
```

```
sftp <file> <remotefile>
```

```
rsync -ravP -e "ssh -i key"  
<localfile> <remotefile>
```

- Users can copy files to the cluster using SCP
- Uses SSH keypair for passwordless transfer
- Data is encrypted in transit
- Many desktop clients available:
 - scp/sftp (Linux, Mac)
 - pscp/psftp (Windows)
 - WinSCP, FileZila (Windows explorer interface)
- Specify keypair, username and path to files:

e.g. `scp -i aws.pub alces@52.50.133.133:/home/alces/file .`

- Tools like **rsync** can make copying faster by skipping files that haven't changed:

e.g. `rsync -ravP -e "ssh -i aws.pub" /storage/data
alces@52.50.133.133:/home/alces/.`

Archiving Data

Useful commands:

```
alces storage enable
alces storage configure
alces storage avail
Alces storage use
alces storage get
alces storage put
```

- The **alces storage** command allows users to store files in S3, dropbox and POSIX services
- Use the **alces storage enable** command:
 - **s3** - Amazon S3 object storage
 - **dropbox** - Dropbox object storage
 - **posix** - An existing POSIX path
- Use the **alces storage configure** command to set up a new storage service

e.g. `alces storage configure mys3 s3`

- Select a service with **avail** and **use** commands
- Upload and download files with **get** and **put**

e.g. `alces storage put localfile s3://remotefile`
`alces storage --recursive get s3://jobs .`
`alces storage mkbucket newbucket`

Running an Application on the login node

Useful commands:

```
alces session start  
xrandr  
module load
```

- Batch jobs are submitted to run on compute nodes via the job scheduler
- Users can also run interactively on the login node
 - Example work-flow:

```
alces session start gnome
```

<connect to VNC desktop>

<from the desktop **Application** menu, select **Terminal**>

```
module load apps/rstudio  
rstudio
```

<inside Rstudio console>

```
demo (graphics)
```

Using the Job Scheduler

- Users submit jobs to a batch queue which are then run on compute nodes
- The job scheduler is responsible for:
 - Starting and stopping jobs
 - Assembling machinefile for parallel jobs
 - Storing printed job output in a file
- Users provide a **job script** which instructs the scheduler what to do to run their application
- Users can use **directives** to instruct the scheduler how to run their job; these may be
 - Included in the user's job-script
 - Provided on the command-line at submission time

Writing a job script

Useful commands:

```
qstat
```

```
qsub <job-script name>
```

```
cat <output file>
```

- A simple job-script:

```
#!/bin/bash -l
echo Hello from job $JOB_ID
sleep 60
echo I was run on node $HOSTNAME
```

- Submit the job using **qsub <jobscript>**
 - View the job status using **qstat**
 - Output shows your unique job ID
 - Job status (r=running, qw=queuing/waiting)
 - By default, job output stored in home-dir
 - Filename is **~/<jobscript-name>.o<job-ID>**
- e.g. `cat ~/myjob.sh.o5`

Job scheduler directives

Useful commands:

```
qstat
```

```
qsub <job-script name>
```

```
cat <output file>
```

- The directive “**-o <filename>**” instructs the scheduler where to store job output
 - The directive can be added at submission time:

E.g. `qsub -o /home/alces/myoutput myjob.sh`

- The directive can also be added in the job-script:

```
#!/bin/bash -l
#$ -o /home/alces/myoutput
echo Hello from job $JOB_ID
sleep 60
echo I was run on node $HOSTNAME
```

- Lines in a job-script starting with `#$` are evaluated at submission time

Job scheduler directives

Useful commands:

`qstat`

`qsub <directive> <job-script name>`

- More scheduler directives:

<code>-o <filename></code>	Set job output file location
<code>-hold <job-ID,...></code>	Wait until the specified job IDs have completed running
<code>-pe <PEname> <slots></code>	Use named parallel environment, requesting a number of slots
<code>-p <priority></code>	Set the job priority (0=normal, -1000=low)
<code>-r <yes no></code>	Set the job to automatically rerun if terminated
<code>-l h_vmem=16G</code>	Request an amount of usable memory per slot (defaults to total / slots)
<code>-l rt=36:00:00</code>	Sets max job runtime to 36 hours

Example: Running an MPI job

Useful commands:

```
qsub imb.sh
```

- The **Intel Messaging Benchmark** (IMB) tests latency and bandwidth between compute nodes
- Write a simple job script that launches the IMB application, and submit to the cluster scheduler
- Watch the output to view results

```
#!/bin/bash -l
#$ -j y -N imb
#$ -o $HOME/imb.$JOB_ID.out
#$ -pe mpinodes-verbose 2 -cwd -V

module load apps/imb
mpirun -np 2 -npernode 1 IMB-MPI1
```

Template job-scripts

Useful commands:

```
alces template list
```

```
alces template copy <name>  
    <new jobscript>
```

```
qsub <jobscript>
```

- A reference set of job-scripts are available using the **alces template** command
 - **alces template list** shows a list of available templates
 - **simple** - single-core serial job
 - **simple-array** - multiple single-core serial jobs
 - **smp** - multi-threaded job on a single node
 - **mpi-slots** - parallel job across multiple nodes, allocated by number of CPU cores
 - **mpi-nodes** - parallel job across multiple nodes, allocated by number of nodes
- The **alces template copy** command allows users to create new job scripts using a template

```
[alces@login1(tiny) jobs]$ alces template copy smp  
mysmpjob.sh  
alces template copy: template 'smp' copied to  
'mysmpjob.sh'
```

Managing Jobs

Useful commands:

`qstat -j <job-ID>`

`qacct -j <job-ID>`

`qsub <job-script name>`

`cat <output file>`

- Users can view the job queue using the **qstat** command; use **qstat -f** for full output
- Individual jobs can be deleted using **qdel <job-ID>**
- Job resource usage is shown using the **qacct <job-ID>** command (once job is completed)
- The status of compute nodes is shown using the **qhost** command

```
[alces-cluster@vlogin1(tiny) jobs]$ qhost
```

HOSTNAME	ARCH	NCPU	LOAD	MEMTOT	MEMUSE	SWAPTO	SWAPUS
-----	-----	-----	-----	-----	-----	-----	-----
global	-	-	-	-	-	-	-
vnnode01	linux-x64	20	0.01	58.8G	620.8M	32.0M	0.0
vnnode02	linux-x64	20	0.02	58.8G	638.0M	32.0M	0.0
vnnode03	linux-x64	20	0.01	58.8G	651.0M	32.0M	0.0
vnnode04	linux-x64	20	0.01	58.8G	622.5M	32.0M	0.0
vnnode05	linux-x64	20	0.01	58.8G	682.0M	32.0M	0.0
vnnode06	linux-x64	20	0.01	58.8G	642.6M	32.0M	0.0
vnnode07	linux-x64	20	0.01	58.8G	667.6M	32.0M	0.0
vnnode08	linux-x64	20	0.01	58.8G	642.5M	32.0M	0.0
vnnode09	linux-x64	20	0.01	58.8G	661.3M	32.0M	0.0



Terminating your cluster

- Before termination, store any job-scripts, data and output files created
 - Use **alces storage put --recursive** command to check-in data you want to keep
- List any applications you've installed
 - Use **module list** to view applications installed
 - Use **alces gridware depot export** to save your applications as a TAR file
- Delete the cluster stack from the *AWS CloudFormation* console
- Once deleted, users are no longer billed for the instances or software



Start computing today

- Many applications to explore
- Lots more help available online
 - <http://alces-flight-appliance-docs.readthedocs.org/en/latest/>



alcesflight

Thank you

www.alces-flight.com

<http://alces-flight.github.io>